

بِسْمِ اللَّهِ الرَّحْمَنِ الرَّحِيمِ



# MongoDB

HOSSEIN FORGHANI

MAKTAB SHARIF



# contents

- ▶ How to design Mongo database?
- ▶ Index
- ▶ Aggregation & pipelines
- ▶ Dump & restore
- ▶ Replication & sharding

# Data Modeling

- ▶ Embedded data model
- ▶ Normalized data model
- ▶ Best practice: mixture of the above!

# Embedded Data Model

```
{
  _id: ,
  Emp_ID: "10025AE336"
  Personal_details:{
    First_Name: "Radhika",
    Last_Name: "Sharma",
    Date_Of_Birth: "1995-09-26"
  },
  Contact: {
    e-mail: "radhika_sharma.123@gmail.com",
    phone: "9848022338"
  },
  Address: {
    city: "Hyderabad",
    Area: "Madapur",
    State: "Telangana"
  }
}
```

# Normalized Data Model

## Employee:

```
{
  _id: <ObjectId101>,
  Emp_ID: "10025AE336"
}
```

## Personal\_details:

```
{
  _id: <ObjectId102>,
  empDocID: " ObjectId101",
  First_Name: "Radhika",
  Last_Name: "Sharma",
  Date_Of_Birth: "1995-09-26"
}
```

## Contact:

```
{
  _id: <ObjectId103>,
  empDocID: " ObjectId101",
  e-mail: "radhika_sharma.123@gmail.com",
  phone: "9848022338"
}
```

## Address:

```
{
  _id: <ObjectId104>,
  empDocID: " ObjectId101",
  city: "Hyderabad",
  Area: "Madapur",
  State: "Telangana"
}
```

# Design Notes

- ▶ Design your schema according to user requirements
- ▶ Combine objects into one document if you will use them together
- ▶ Since MongoDB does not have JOIN, make sure you do not need join between collections!
- ▶ Optimize your schema for most frequent use cases
- ▶ Pre-compute complex aggregations

# Indexing

- ▶ Indexes are special data structures, that store a small portion of the data set in an easy-to-traverse form
- ▶ To create an index on the field "KEY" in an ascending order:

```
db.COLLECTION_NAME.createIndex({KEY:1})
```

- ▶ Also multiple fields:

```
db.mycol.createIndex({"title":1,"description":-1})
```

- ▶ This index can support a sort on {"title":1,"description":-1}
- ▶ createIndex has some options: background, unique, name, ...



# Indexing – cont.

- ▶ To drop index:

```
db.COLLECTION_NAME.dropIndex({KEY:1})
```

- ▶ See also:
  - ▶ createIndexes
  - ▶ dropIndexes
  - ▶ getIndexes

# Aggregation

- ▶ Aggregations operations process data records and return computed results
- ▶ Equivalent to SQL “group by” and aggregation functions

```
db.COLLECTION_NAME.aggregate(pipeline, options)
```

In a collection you have the following data:

```
{
  _id: ObjectId(7df78ad8902c)
  title: 'MongoDB Overview',
  description: 'MongoDB is no sql database',
  by_user: 'tutorials point',
  url: 'http://www.tutorialspoint.com',
  tags: ['mongodb', 'database', 'NoSQL'],
  likes: 100
},
{
  _id: ObjectId(7df78ad8902d)
  title: 'NoSQL Overview',
  description: 'No sql database is very fast',
  by_user: 'tutorials point',
  url: 'http://www.tutorialspoint.com',
  tags: ['mongodb', 'database', 'NoSQL'],
  likes: 10
},
{
  _id: ObjectId(7df78ad8902e)
  title: 'Neo4j Overview',
  description: 'Neo4j is no sql database',
  by_user: 'Neo4j',
  url: 'http://www.neo4j.com',
  tags: ['neo4j', 'database', 'NoSQL'],
  likes: 750
},
```

## Aggregation example:

Counts the number of “by\_user” values

pipelines

```
➤ db.mycol.aggregate(  
  [  
    { $group : {  
      _id : "$by_user",  
      num_tutorial : { $sum : 1 }  
    }  
  ] )  
  
{ "_id" : "tutorials point", "num_tutorial" : 2 }  
{ "_id" : "Neo4j", "num_tutorial" : 1 }
```

Equivalent to SQL group by

Equivalent to SQL  
aggregation function

# Aggregation – cont.

- ▶ Sums up “likes” values for each “by\_user” value

```
db.mycol.aggregate([{$group : {  
  _id : "$by_user",  
  likes : {$sum : "$likes"}  
}}])
```

- ▶ Also \$avg, \$min, \$max

# Aggregation Pipeline

- ▶ Aggregation steps
  - ▶ **\$project**: same as find() project
  - ▶ **\$match**: same as find() filter
  - ▶ **\$group**: grouping and aggregation
  - ▶ **\$sort**: same as sort()
  - ▶ **\$skip**: same as skip()
  - ▶ **\$limit**: same as limit()
  - ▶ **\$unwind**: unpack the arrays
- ▶ All are optional



# Aggregation Pipeline – cont.

- ▶ Selects documents with status equal to "A"
- ▶ Groups them by the `cust_id` field and calculates the total for each `cust_id` field from the `sum` of the `amount` field
- ▶ And sorts the results by the `total` field in descending order

```
db.orders.aggregate([  
    { $match: { status: "A" } },  
    { $group: { _id: "$cust_id", total: { $sum: "$amount" } } },  
    { $sort: { total: -1 } }  
])
```

# Aggregation Pipeline – cont.

- ▶ To handle large datasets, set `allowDiskUse` option to true to enable writing data to temporary files

```
db.stocks.aggregate(  
  [  
    { $project : { cusip: 1, date: 1, price: 1, _id: 0 } },  
    { $sort : { cusip : 1, date: 1 } }  
  ],  
  { allowDiskUse: true }  
)
```



# Some tips

- ▶ Set group null if there is no grouping:

```
db.orders.aggregate([ { $group: { _id: null, total: { $sum: "$amount" } } } ])
```

- ▶ You have a wide range of aggregation operators:

\$sum	\$first	\$dayOfWeek	\$size	\$add	\$toInt
\$max	\$last	\$dayOfYear	\$floor	\$subtract	\$toDecimal
\$min	\$hour	\$sin	\$ceil	\$divide	\$toLong
\$avg	\$dayOfMonth	\$cos	\$multiply	\$toDate	\$toString

- ▶ For more reading:

- ▶ <https://docs.mongodb.com/manual/reference/operator/aggregation-pipeline/>

# \$unwind

- ▶ Deconstructs an array field from the input documents to output a document for each element

```
{ $unwind: <field path> }
```

- ▶ Or:

```
{  
  $unwind:  
  {  
    path: <field path>,  
    includeArrayIndex: <string>,  
    preserveNullAndEmptyArrays: <boolean>  
  }  
}
```

# \$unwind Example

► Consider this collection:

```
{ "_id" : 1, "item" : "ABC", price: NumberDecimal("80"), "sizes": [ "S", "M", "L"] }  
{ "_id" : 2, "item" : "EFG", price: NumberDecimal("120"), "sizes" : [ ] }  
{ "_id" : 3, "item" : "IJK", price: NumberDecimal("160"), "sizes": "M" }  
{ "_id" : 4, "item" : "LMN", price: NumberDecimal("10") }  
{ "_id" : 5, "item" : "XYZ", price: NumberDecimal("5.75"), "sizes" : null }
```

## \$unwind Example – cont.

```
db.inventory2.aggregate( [  
    // First Stage  
    { $unwind: { path: "$sizes", preserveNullAndEmptyArrays: true } },  
    // Second Stage  
    { $group:  
        { _id: "$sizes", averagePrice: { $avg: "$price" } }  
    },  
    // Third Stage  
    { $sort: { "averagePrice": -1 } }  
] )
```

# \$unwind Example – cont.

► Stage 1: \$unwind

```
{ "_id" : 1, "item" : "ABC", "price" : NumberDecimal("80"), "sizes" : "S" }  
{ "_id" : 1, "item" : "ABC", "price" : NumberDecimal("80"), "sizes" : "M" }  
{ "_id" : 1, "item" : "ABC", "price" : NumberDecimal("80"), "sizes" : "L" }  
{ "_id" : 2, "item" : "EFG", "price" : NumberDecimal("120") }  
{ "_id" : 3, "item" : "IJK", "price" : NumberDecimal("160"), "sizes" : "M" }  
{ "_id" : 4, "item" : "LMN", "price" : NumberDecimal("10") }  
{ "_id" : 5, "item" : "XYZ", "price" : NumberDecimal("5.75"), "sizes" : null }
```

# \$unwind Example – cont.

► Stage 2: \$group

```
{ "_id" : "S", "averagePrice" : NumberDecimal("80") }  
{ "_id" : "L", "averagePrice" : NumberDecimal("80") }  
{ "_id" : "M", "averagePrice" : NumberDecimal("120") }  
{ "_id" : null, "averagePrice" : NumberDecimal("45.25") }
```

► Stage 3: \$sort

```
{ "_id" : "M", "averagePrice" : NumberDecimal("120") }  
{ "_id" : "L", "averagePrice" : NumberDecimal("80") }  
{ "_id" : "S", "averagePrice" : NumberDecimal("80") }  
{ "_id" : null, "averagePrice" : NumberDecimal("45.25") }
```

# Dump and Restore

```
mongodump --host="mongodb0.example.com:27017" --port=27017 --db=<db_name> --out=<pat
```

```
mongorestore --host="mongodb0.example.com:27017" --port=27017 --db=<db_name> <path>
```

- ▶ For more options refer to the references
- ▶ See Also:
  - ▶ mongoexport
  - ▶ mongoimport

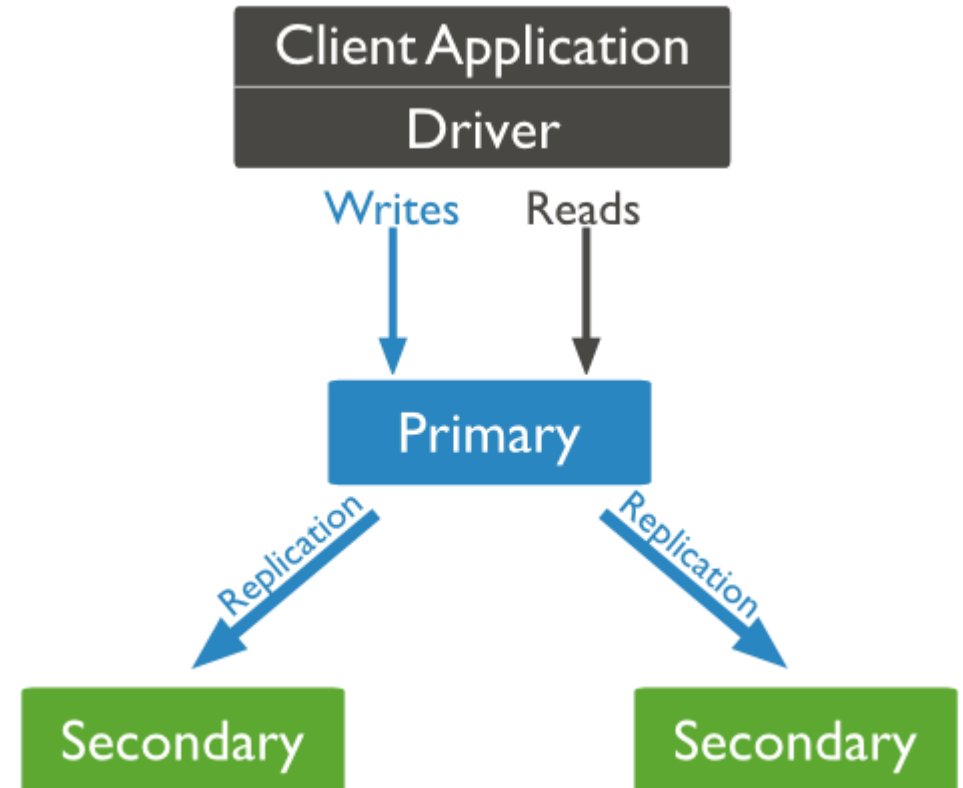
# Replication

- ▶ Replication is the process of **synchronizing data across multiple servers**
- ▶ Cost
  - ▶ Provides **redundancy**
- ▶ Benefit
  - ▶ Increases data **availability**
  - ▶ Keeps your data **safe**
  - ▶ **Disaster recovery**
  - ▶ **No downtime** for maintenance



# Replication – cont.

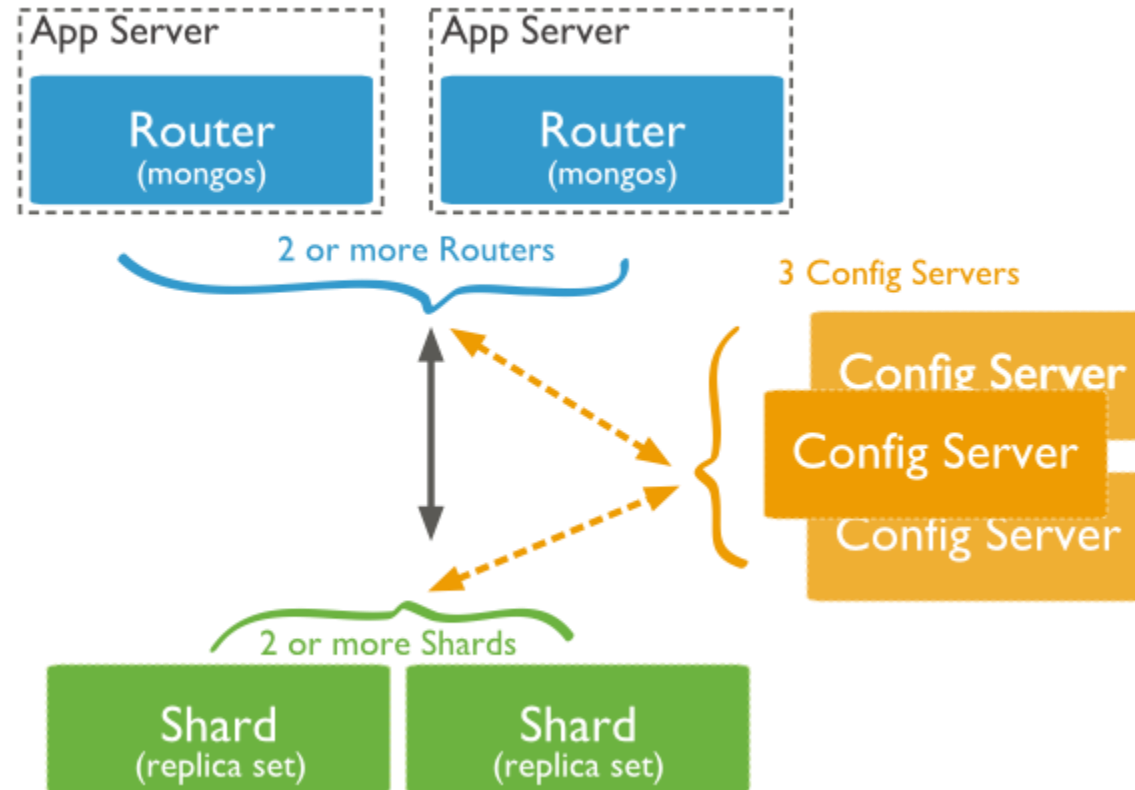
- ▶ Replica set:
  - ▶ 2 or more nodes
  - ▶ A **primary** node and many **secondary** nodes
  - ▶ At the time of automatic failover or maintenance, **election** establishes for primary and a new primary node is elected
  - ▶ After the recovery of failed node, it again join the replica set and works as a **secondary node**
- ▶ To see how to setup a replica set refer to the references



# Sharding

- ▶ Sharding is the process of storing data records **across multiple machines** on a **single server** to handle data growth
- ▶ Sharding vs Replication:
  - ▶ **Benefit**: no need to buy any server or add any RAM, CPU, Disk
  - ▶ **Cost**: increased complexity in infrastructure and maintenance

# Sharding – cont.



# References

- ▶ <https://www.tutorialspoint.com/mongodb>
- ▶ <https://docs.mongodb.com/manual/reference>

# Any Question?